

Lab2

Threads

Name	
ID	Sohayla Mohammed
	32

Matrix Multiplication

Problem Statement

It is required to implement two variations of matrix multiplication:

- The computation of each element of the output matrix happens in a thread.
- The computation of each row of the output matrix happens in a thread.

For both variations, it's required to compute the elapsed time for each of them and compare them.

Objectives

- Understanding threads concepts.
-

Design and Implementation

Tools & Libraries Used

- Coded with c++.
 - Libraries
-

-
- pthread
-

Data Structures

- vector<int>
 - Holds the size of the input matrix.
 - vector<vector<int>>
 - Holds the input matrices .
-

Built in

- Strings
 - getline()
 - Threads
 - pthread_create(), pthread_join()
 - Synchronization
 - mutex
-

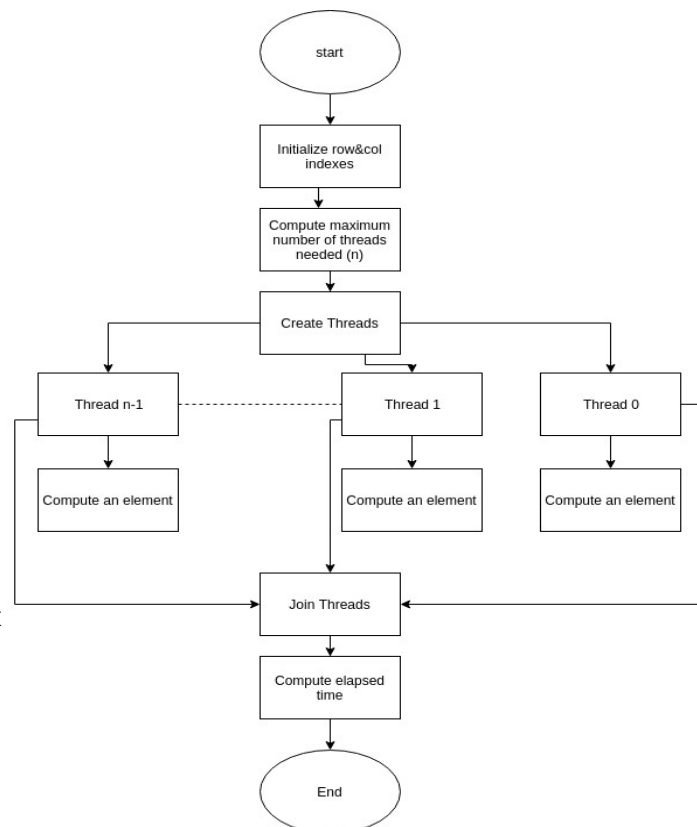
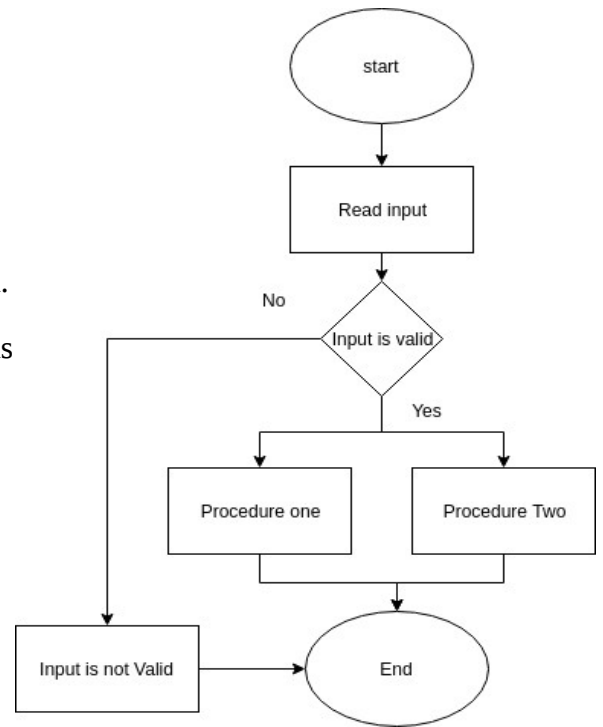
Assumptions

- The input is given in the format shown in the “matrix-readme.txt” file.
- The output file contains
 - result matrices from the two procedures.
 - The elapsed time of each procedure.

Functionality

Overall

- Upon starting the program:
 - Reads data from file “input.txt”
 - If file does not exist the program is terminated.
 - If the format of the file is wrong the program is terminated
 - If the input is not valid for matrix multiplication, the program is terminated.
 - When the input is valid
 - first call is to procedure one which will compute each element in a thread.
 - Second call is to procedure two which will compute each row in a thread.
 - The output matrix and elapsed time from each procedure is written to “output.txt” file.
- In the first procedure :
 - Row and Column indices are initialized to zero.
 - Max number of threads is the multiplication of the number of rows in first matrix and the number of columns in the second matrix.
 - An array of threads is created.
 - Each thread is created to compute an element in the matrix.
 - All threads are joined later.



-
- An elapsed time is calculated from the beginning until the last thread is joined.
 - In the second procedure :
 - The same as first but
 - The maximum number of threads is equal to the number of rows in the first matrix.
 - Each thread computes a row in the matrix.

Functions

- **void** begin(string filename);
 - Begins the execution of the program.
 - Parameters :
 - string filename denotes the name of the input file.
 - Return type : void
- **void** readFileMult(string filename);
 - Reads data from input file to datastructures.
 - Parameters :
 - string filename denotes the name of the input file
- Two Helper functions :
 - Act as a bridge between pthread_create() and the needed function.
- **void** matrixMult1();
 - Represents procedure one, which creates a thread for each element computed.
 - Parameters None.
 - Return type void.
- **void** matrixMult2();
 - Represents procedure twp, which creates a thread for each row computed.
 - Parameters None.
 - Return type void.
- **void*** computeAnElement(**void**);
 - Computes each element from multiplying the two matrices.
 - Parameters NONE.
 - Return type *void (auto)
- **void*** computeARow(**void**);
 - Computes each row from multiplying the two matrices.
 - Parameters NONE.

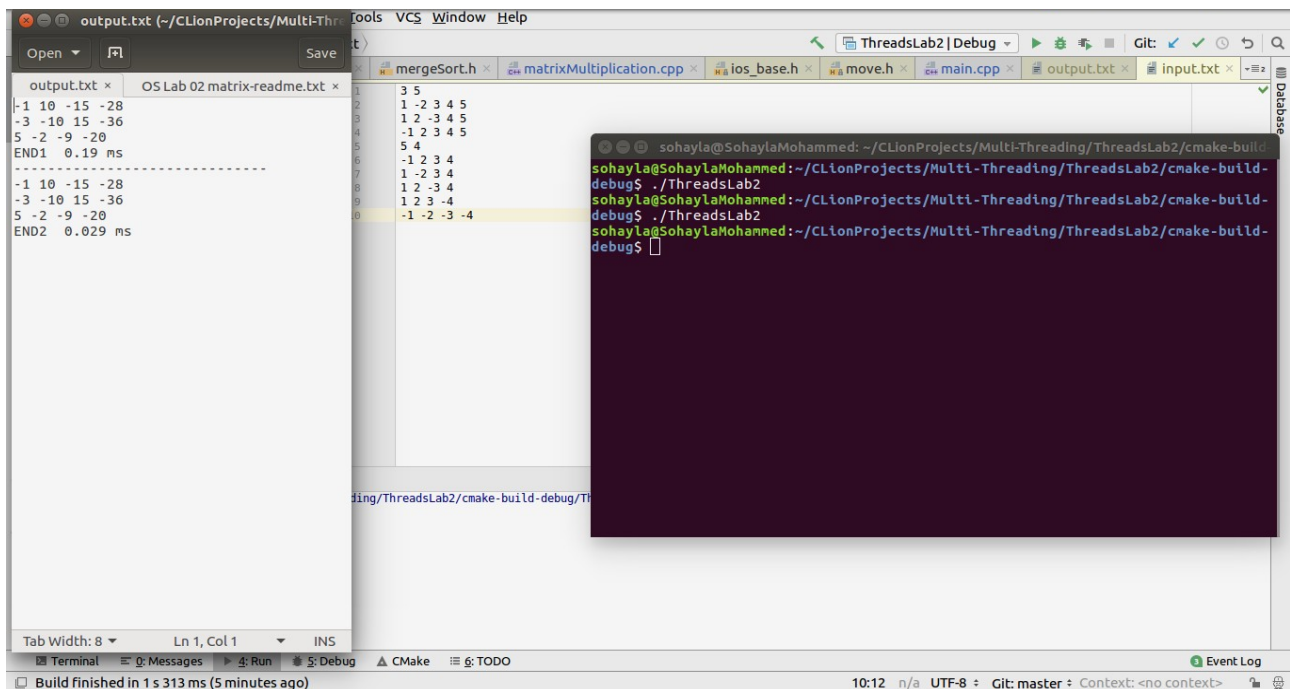
- Return type *void (auto)
- **void** writeOutputFile();
 - Writes the results into “output.txt” file.
 - Parameters None.
 - Return type void.

Sample Runs

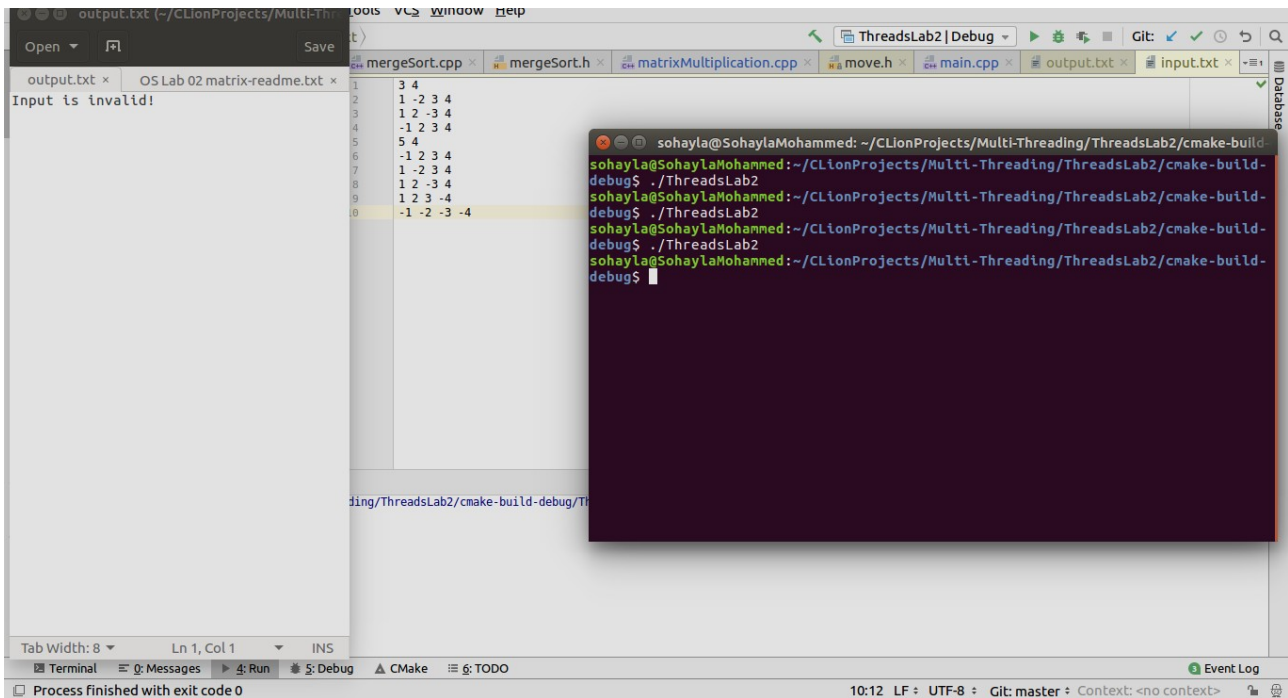
Input file does not exist, output file :



Given input :



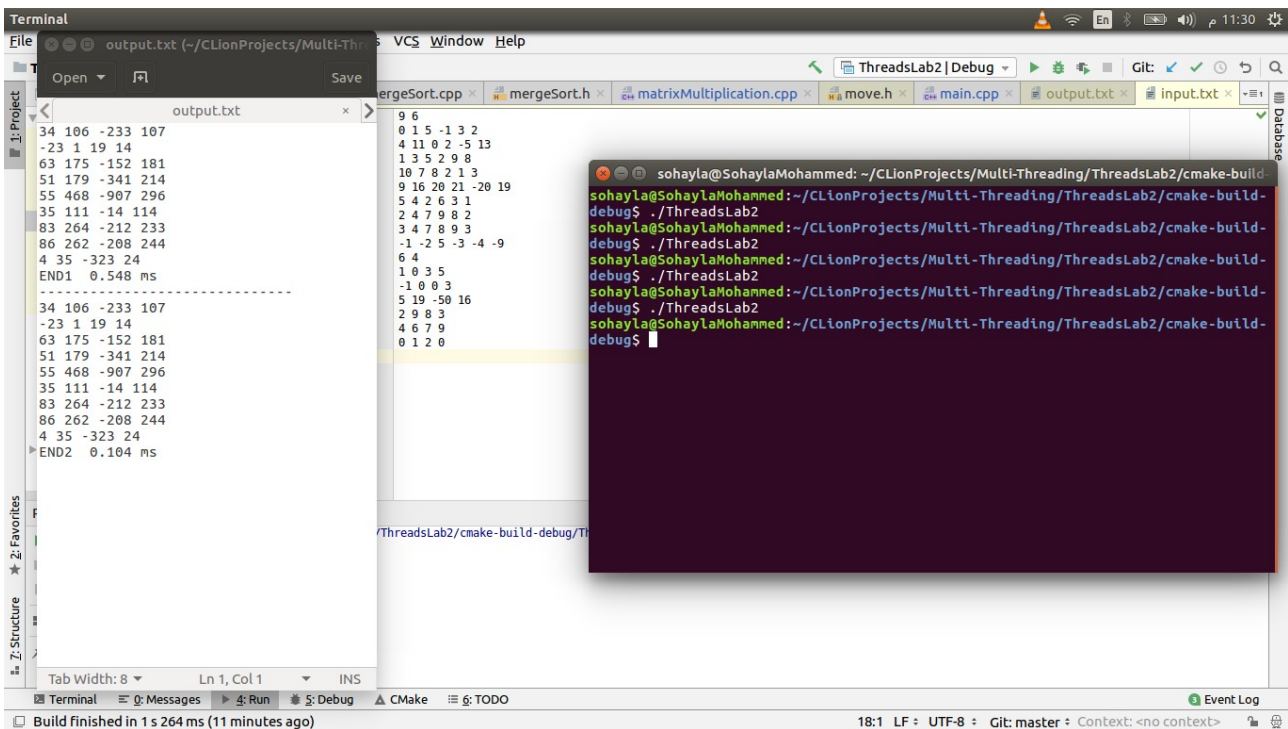
Invalid input



The screenshot shows the CLion IDE interface. The main editor displays a C++ file named `mergeSort.cpp` with a matrix of numbers. The `main.cpp` file is also visible, showing a loop that reads input from `input.txt`. The `output.txt` file shows the output: "Input is invalid!". A terminal window is open, showing the command prompt and the execution of the program.

```
sohayla@SohaylaMohammed: ~/CLionProjects/Multi-Threading/ThreadsLab2/cmake-build-debug$ ./ThreadsLab2
sohayla@SohaylaMohammed:~/CLionProjects/Multi-Threading/ThreadsLab2/cmake-build-debug$ ./ThreadsLab2
sohayla@SohaylaMohammed:~/CLionProjects/Multi-Threading/ThreadsLab2/cmake-build-debug$ ./ThreadsLab2
sohayla@SohaylaMohammed:~/CLionProjects/Multi-Threading/ThreadsLab2/cmake-build-debug$
```

Given input



The screenshot shows the CLion IDE interface. The main editor displays a C++ file named `mergeSort.cpp` with a matrix of numbers. The `main.cpp` file is also visible, showing a loop that reads input from `input.txt`. The `output.txt` file shows the output: a matrix of numbers. A terminal window is open, showing the command prompt and the execution of the program.

```
sohayla@SohaylaMohammed: ~/CLionProjects/Multi-Threading/ThreadsLab2/cmake-build-debug$ ./ThreadsLab2
sohayla@SohaylaMohammed:~/CLionProjects/Multi-Threading/ThreadsLab2/cmake-build-debug$ ./ThreadsLab2
sohayla@SohaylaMohammed:~/CLionProjects/Multi-Threading/ThreadsLab2/cmake-build-debug$ ./ThreadsLab2
sohayla@SohaylaMohammed:~/CLionProjects/Multi-Threading/ThreadsLab2/cmake-build-debug$ ./ThreadsLab2
sohayla@SohaylaMohammed:~/CLionProjects/Multi-Threading/ThreadsLab2/cmake-build-debug$ ./ThreadsLab2
sohayla@SohaylaMohammed:~/CLionProjects/Multi-Threading/ThreadsLab2/cmake-build-debug$
```

Time Elapse

	Procedure one	Procedure Two
Test 1	0.19	0.029
Test 2	0.548	0.104
Test 3	1.084	0.26
Test 4	0.772	0.166
Test 5	0.424	0.035

Conclusion

- Procedure two takes less time than procedure one.
- As procedure one creates more threads, as each thread creation and termination takes more overhead time.
- Number of threads in the first procedure is equals ($m*n$) and in the second procedure equals (m).
- If the number of columns in the second matrix is 1, the difference between two procedures time is less.

Merge Sort

Problem Statement

Merge sort is an $O(n \log n)$ comparison-based sorting algorithm. It is a divide and conquer algorithm.

It's required to implement it using Pthreads. Each time the list is divided; two threads are created to do merge-sort on each half separately. This step is repeated recursively until each sub-list has only one element.

Objectives

- Understanding threads concepts.

Design and Implementation

Tools & Libraries Used

- Coded with c++.
- Libraries
 - pthread

Data Structures

- `vector<int>`
 - Holds the input array.
 - Holds the output array.
- `vector<pthread_t>`

-
- Holds the threads created
-

Built in

- Strings
 - `getline()`
 - Threads
 - `pthread_create()`, `pthread_join()`
 - Synchronization
 - `mutex`
-

Assumptions

- The input is given in the format shown in the “merge-readme.txt” file.
 - The output is printed after the finishing of the program.
-

Functionality

Overall

- Upon starting the program:
 - Reads data from file “input.txt”
 - If file does not exist the program is terminated.
 - If the format of the file is wrong the program is terminated
 - When the input is valid
 - A thread is created as the parent thread.
 - Then recursively :

-
- Each array in a thread is divided into 2 until there exist a thread with only one element
 - Then each thread sort it's array then merged it with it's sibling thread
 - This procedure is done from bottom up.
 - All the threads join.
 - The output array is printed after it's done.

Functions

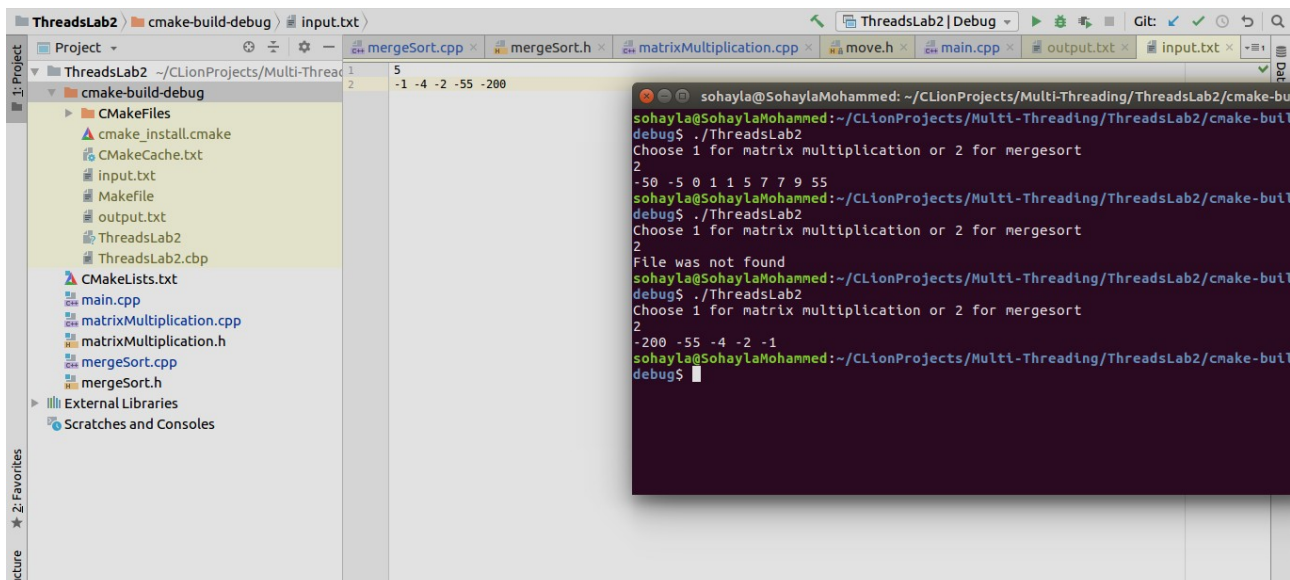
- **void** begin(string filename);
 - Begins the execution of the program.
 - Parameters :
 - string filename denotes the name of the input file.
 - Return type : void
- **void** readFileMerge(string filename);
 - Reads data from input file to datastructures.
 - Parameters :
 - string filename denotes the name of the input file
 - Return type void
- one Helper function :
 - Acts as a bridge between pthread_create() and the needed function.
- **void** mainThread();
 - Creates the parent thread which in return responsible for all the other ones.
 - Parameters NONE
 - Return type void
- **void * mergeSortF(void* arg);**
 - Divides the given array "arg" into two arrays, then creates two threads to merge sort each one.
 - Parameters :
 - arg represents the given sub-array in current recursive call.
 - Return type void
- **void merge(vector<int>* a, vector<int> l, vector<int> r);**
 - Merges two sorted arrays l & r into a.
 - Parameters :
 - l → left sorted part of the array
 - r → right sorted part of the array.
 - a → holds the merging result of l & r.

Sample Runs

No input File

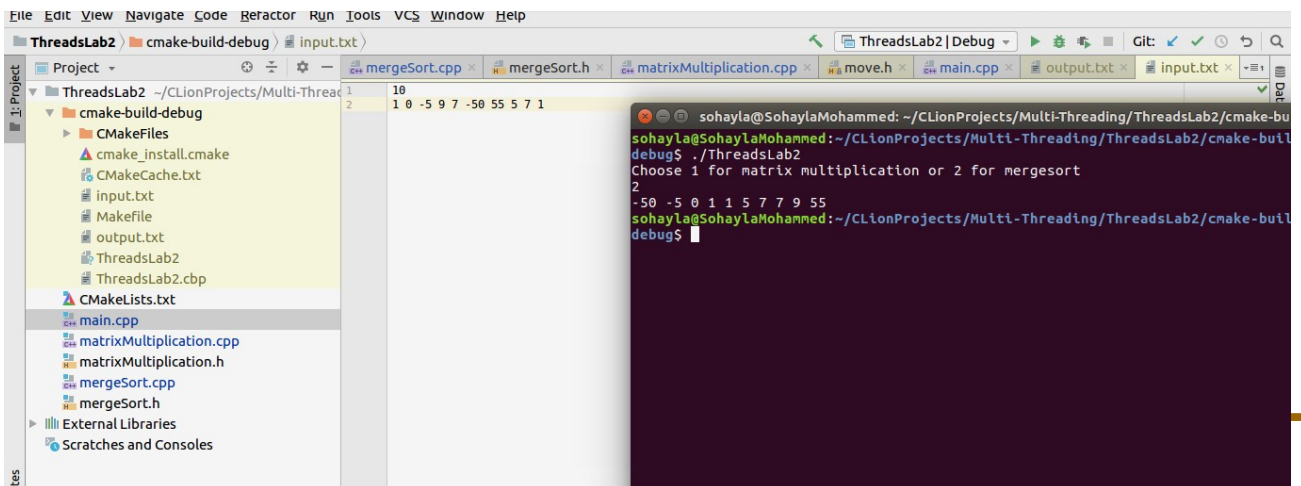
```
sohayla@SohaylaMohammed: ~/CLionProjects/Multi-Threading/ThreadsLab2/cmake-build-debug
sohayla@SohaylaMohammed:~/CLionProjects/Multi-Threading/ThreadsLab2/cmake-build-debug$ ./ThreadsLab2
Choose 1 for matrix multiplication or 2 for mergesort
2
-50 -5 0 1 1 5 7 7 9 55
sohayla@SohaylaMohammed:~/CLionProjects/Multi-Threading/ThreadsLab2/cmake-build-debug$ ./ThreadsLab2
Choose 1 for matrix multiplication or 2 for mergesort
2
File was not found
sohayla@SohaylaMohammed:~/CLionProjects/Multi-Threading/ThreadsLab2/cmake-build-debug$
```

Given input



```
ThreadsLab2 | Debug
sohayla@SohaylaMohammed:~/CLionProjects/Multi-Threading/ThreadsLab2/cmake-build-debug$ ./ThreadsLab2
Choose 1 for matrix multiplication or 2 for mergesort
2
-50 -5 0 1 1 5 7 7 9 55
sohayla@SohaylaMohammed:~/CLionProjects/Multi-Threading/ThreadsLab2/cmake-build-debug$ ./ThreadsLab2
Choose 1 for matrix multiplication or 2 for mergesort
2
File was not found
sohayla@SohaylaMohammed:~/CLionProjects/Multi-Threading/ThreadsLab2/cmake-build-debug$ ./ThreadsLab2
Choose 1 for matrix multiplication or 2 for mergesort
2
-200 -55 -4 -2 -1
sohayla@SohaylaMohammed:~/CLionProjects/Multi-Threading/ThreadsLab2/cmake-build-debug$
```

Given input



```
ThreadsLab2 | Debug
sohayla@SohaylaMohammed:~/CLionProjects/Multi-Threading/ThreadsLab2/cmake-build-debug$ ./ThreadsLab2
Choose 1 for matrix multiplication or 2 for mergesort
2
-50 -5 0 1 1 5 7 7 9 55
sohayla@SohaylaMohammed:~/CLionProjects/Multi-Threading/ThreadsLab2/cmake-build-debug$ ./ThreadsLab2
Choose 1 for matrix multiplication or 2 for mergesort
2
-200 -55 -4 -2 -1
sohayla@SohaylaMohammed:~/CLionProjects/Multi-Threading/ThreadsLab2/cmake-build-debug$
```